

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats.stats import pearsonr
import numpy as np
from sklearn.model_selection import train_test_split
import seaborn as sns
```

```
In [2]: raw_data= pd.read_csv('Bank_Personal_Loan_Modelling.csv')
```

```
In [3]: raw_data.head()
```

Out[3]:

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	Personal Loan	Secu Ac
0	1	25	1	49	91107	4	1.6	1	0	0	
1	2	45	19	34	90089	3	1.5	1	0	0	
2	3	39	15	11	94720	1	1.0	1	0	0	
3	4	35	9	100	94112	1	2.7	2	0	0	
4	5	35	8	45	91330	4	1.0	2	0	0	

```
In [4]: np.log(raw_data)
```

```
C:\Users\ppragallapati\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel_launcher.py:1: RuntimeWarning: divide by zero encountered in log
    """Entry point for launching an IPython kernel.
C:\Users\ppragallapati\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel_launcher.py:1: RuntimeWarning: invalid value encountered in log
    """Entry point for launching an IPython kernel.
```

In [5]: `raw_data.corr()`

Out[5]:

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	E
<b>ID</b>	1.000000	-0.008473	-0.008326	-0.017695	0.013432	-0.016797	-0.024675	
<b>Age</b>	-0.008473	1.000000	0.994215	-0.055269	-0.029216	-0.046418	-0.052012	
<b>Experience</b>	-0.008326	0.994215	1.000000	-0.046574	-0.028626	-0.052563	-0.050077	
<b>Income</b>	-0.017695	-0.055269	-0.046574	1.000000	-0.016410	-0.157501	0.645984	
<b>ZIP Code</b>	0.013432	-0.029216	-0.028626	-0.016410	1.000000	0.011778	-0.004061	
<b>Family</b>	-0.016797	-0.046418	-0.052563	-0.157501	0.011778	1.000000	-0.109275	
<b>CCAvg</b>	-0.024675	-0.052012	-0.050077	0.645984	-0.004061	-0.109275	1.000000	
<b>Education</b>	0.021463	0.041334	0.013152	-0.187524	-0.017377	0.064929	-0.136124	
<b>Mortgage</b>	-0.013920	-0.012539	-0.010582	0.206806	0.007383	-0.020445	0.109905	
<b>Personal Loan</b>	-0.024801	-0.007726	-0.007413	0.502462	0.000107	0.061367	0.366889	

## Checking Outliers

In [6]: `raw_data.describe()`

Out[6]:

	ID	Age	Experience	Income	ZIP Code	Family	C
<b>count</b>	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.0
<b>mean</b>	2500.500000	45.338400	20.104600	73.774200	93152.503000	2.396400	1.9
<b>std</b>	1443.520003	11.463166	11.467954	46.033729	2121.852197	1.147663	1.7
<b>min</b>	1.000000	23.000000	-3.000000	8.000000	9307.000000	1.000000	0.0
<b>25%</b>	1250.750000	35.000000	10.000000	39.000000	91911.000000	1.000000	0.7
<b>50%</b>	2500.500000	45.000000	20.000000	64.000000	93437.000000	2.000000	1.5
<b>75%</b>	3750.250000	55.000000	30.000000	98.000000	94608.000000	3.000000	2.5
<b>max</b>	5000.000000	67.000000	43.000000	224.000000	96651.000000	4.000000	10.0


```
In [7]: sns.boxplot(data=raw_data['Experience'])
Experience= raw_data['Experience']
count=0
for i in range(5000):
    if Experience[i]< 0:
        count+=1
        print(Experience[i])
        Experience[i]=abs(Experience[i])
        #print(Experience[i])
print(count)
#Experience can not be negative. These are outliers. Converted these negative
```

C:\Users\ppragallapati\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel\_launcher.py:8: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
In [8]: #After checking the distributions of all attributes, only experience has wrong
```

```
In [9]: from scipy.stats import norm
plt.bar(raw_data['Personal Loan'], norm.pdf(raw_data['Personal Loan'], 56, 16))
plt.show()
```



In [ ]: `#There are no outlier/ false data in the target variable`

In [10]: `from sklearn.feature_extraction.text import TfidfVectorizer  
final_data = TfidfVectorizer(raw_data)  
final_data= raw_data.drop(['ZIP Code', 'ID'], axis= 1)  
final_data.head()`

Out[10]:

	Age	Experience	Income	Family	CCAvg	Education	Mortgage	Personal Loan	Securities Account	Acco
0	25	1	49	4	1.6	1	0	0	1	
1	45	19	34	3	1.5	1	0	0	1	
2	39	15	11	1	1.0	1	0	0	0	
3	35	9	100	1	2.7	2	0	0	0	
4	35	8	45	4	1.0	2	0	0	0	

## SPLIT DATA 70:30

In [11]: `x_train, x_test, y_train, y_test= train_test_split(final_data.drop(['Personal`

## Logistic Regression

In [12]: `from sklearn.linear_model import LogisticRegression  
logistic= LogisticRegression()  
logistic.fit(x_train, y_train)  
logistic_prediction= logistic.predict(x_test)`

C:\Users\ppragallapati\AppData\Local\Continuum\anaconda3\lib\site-packages  
\sklearn\linear\_model\logistic.py:433: FutureWarning: Default solver will be  
changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.  
FutureWarning)

```
In [13]: ▶ from sklearn.metrics import classification_report
from sklearn import metrics
print(classification_report(y_test, logistic_prediction))
print("Accuracy:", metrics.accuracy_score(y_test, logistic_prediction))
metrics.confusion_matrix(y_test, logistic_prediction)
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	1355
1	0.83	0.62	0.71	145
micro avg	0.95	0.95	0.95	1500
macro avg	0.90	0.80	0.84	1500
weighted avg	0.95	0.95	0.95	1500

Accuracy: 0.9513333333333334

```
Out[13]: array([[1337,  18],
                [ 55,  90]], dtype=int64)
```

## K- NN

```
In [14]: ▶ from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
knn.fit(x_train, y_train)
knn_prediction= knn.predict(x_test)
```

```
In [15]: ▶ print(classification_report(y_test, knn_prediction))
print("Accuracy:", metrics.accuracy_score(y_test, knn_prediction))
metrics.confusion_matrix(y_test, knn_prediction)
```

	precision	recall	f1-score	support
0	0.93	0.97	0.95	1355
1	0.53	0.30	0.38	145
micro avg	0.91	0.91	0.91	1500
macro avg	0.73	0.63	0.67	1500
weighted avg	0.89	0.91	0.89	1500

Accuracy: 0.9066666666666666

```
Out[15]: array([[1317,  38],
                [102,  43]], dtype=int64)
```

## Naive Bayes

```
In [16]: from sklearn.naive_bayes import BernoulliNB
nb= BernoulliNB()
nb.fit(x_train, y_train)
nb_prediction= nb.predict(x_test)
```

```
In [17]: print(classification_report(y_test,nb_prediction))
print("Accuracy:",metrics.accuracy_score(y_test, nb_prediction))
metrics.confusion_matrix(y_test, nb_prediction)
```

	precision	recall	f1-score	support
0	0.91	0.99	0.95	1355
1	0.39	0.06	0.11	145
micro avg	0.90	0.90	0.90	1500
macro avg	0.65	0.53	0.53	1500
weighted avg	0.86	0.90	0.87	1500

Accuracy: 0.9

```
Out[17]: array([[1341, 14],
               [ 136, 9]], dtype=int64)
```

## Confusion Matrices

```
In [18]: metrics.confusion_matrix(y_test, logistic_prediction)
```

```
Out[18]: array([[1337, 18],
               [ 55, 90]], dtype=int64)
```

```
In [19]: metrics.confusion_matrix(y_test, knn_prediction)
```

```
Out[19]: array([[1317, 38],
               [ 102, 43]], dtype=int64)
```

```
In [20]: metrics.confusion_matrix(y_test, nb_prediction)
```

```
Out[20]: array([[1341, 14],
               [ 136, 9]], dtype=int64)
```

## Conclusion

```
In [ ]: #Logistic regression fetches the best results with an accuracy of 95.1%. This
```

