

## Importing Data

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sn
```

```
In [2]: raw_data= pd.read_csv("vehicle-1.csv")
raw_data.head()
```

```
Out[2]:
```

compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_aspect_ra
95	48.0	83.0	178.0	72.0	
91	41.0	84.0	141.0	57.0	
104	50.0	106.0	209.0	66.0	
93	41.0	82.0	159.0	63.0	
85	44.0	70.0	205.0	103.0	

## Data Preprocessing ¶

In [3]: `raw_data.corr()`

Out[3]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio
compactness	1.000000	0.689786	0.791707	0.691081	0.091779
circularity	0.689786	1.000000	0.797180	0.625051	0.154283
distance_circularity	0.791707	0.797180	1.000000	0.771748	0.158684
radius_ratio	0.691081	0.625051	0.771748	1.000000	0.665363
pr.axis_aspect_ratio	0.091779	0.154283	0.158684	0.665363	1.000000
max.length_aspect_ratio	0.148249	0.251407	0.264621	0.450486	0.640000
scatter_ratio	0.812770	0.858265	0.907949	0.738480	0.100000
elongatedness	-0.788736	-0.827246	-0.913020	-0.792946	-0.180000
pr.axis_rectangularity	0.814248	0.856603	0.896273	0.712744	0.070000
max.length_rectangularity	0.676143	0.965729	0.775149	0.571083	0.120000
scaled_variance	0.764361	0.806791	0.865710	0.798294	0.270000
scaled_variance.1	0.818674	0.850863	0.890541	0.725598	0.080000
scaled_radius_of_gyration	0.585845	0.935950	0.706950	0.541325	0.120000
scaled_radius_of_gyration.1	-0.250603	0.053080	-0.227001	-0.181520	0.150000
skewness_about	0.236685	0.144968	0.114665	0.049112	-0.050000
skewness_about.1	0.157670	-0.011869	0.266049	0.174469	-0.030000
skewness_about.2	0.298528	-0.106339	0.146027	0.382912	0.240000
hollows_ratio	0.365552	0.045652	0.333648	0.472339	0.260000


In [4]: `raw_data.describe()`

Out[4]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_aspect_ratio
count	846.000000	841.000000	842.000000	840.000000	844.000000	844.000000
mean	93.678487	44.828775	82.110451	168.888095	61.678910	61.678910
std	8.234474	6.152172	15.778292	33.520198	7.891463	7.891463
min	73.000000	33.000000	40.000000	104.000000	47.000000	47.000000
25%	87.000000	40.000000	70.000000	141.000000	57.000000	57.000000
50%	93.000000	44.000000	80.000000	167.000000	61.000000	61.000000
75%	100.000000	49.000000	98.000000	195.000000	65.000000	65.000000
max	119.000000	59.000000	112.000000	333.000000	138.000000	138.000000

```
In [5]: sn.pairplot(raw_data, diag_kind = 'kde')
plt.show()
```

```
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\nonparametric\kde.py:44
7: RuntimeWarning: invalid value encountered in greater
  X = X[np.logical_and(X > clip[0], X < clip[1])] # won't work for two columns.
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\nonparametric\kde.py:44
7: RuntimeWarning: invalid value encountered in less
  X = X[np.logical_and(X > clip[0], X < clip[1])] # won't work for two columns.
```



```
In [6]: raw_data['class'] = raw_data['class'].replace({'van': 1, 'bus': 2, 'car': 3})
```

In [7]: `raw_data.describe()`

Out[7]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_
count	846.000000	841.000000	842.000000	840.000000	844.000000	
mean	93.678487	44.828775	82.110451	168.888095	61.678910	
std	8.234474	6.152172	15.778292	33.520198	7.891463	
min	73.000000	33.000000	40.000000	104.000000	47.000000	
25%	87.000000	40.000000	70.000000	141.000000	57.000000	
50%	93.000000	44.000000	80.000000	167.000000	61.000000	
75%	100.000000	49.000000	98.000000	195.000000	65.000000	
max	119.000000	59.000000	112.000000	333.000000	138.000000	

In [8]: `data= raw_data.drop(['class'], axis= 1)`  
`data.median()`

Out[8]:

compactness	93.0
circularity	44.0
distance_circularity	80.0
radius_ratio	167.0
pr.axis_aspect_ratio	61.0
max.length_aspect_ratio	8.0
scatter_ratio	157.0
elongatedness	43.0
pr.axis_rectangularity	20.0
max.length_rectangularity	146.0
scaled_variance	179.0
scaled_variance.1	363.5
scaled_radius_of_gyration	173.5
scaled_radius_of_gyration.1	71.5
skewness_about	6.0
skewness_about.1	11.0
skewness_about.2	188.0
hollows_ratio	197.0
dtype:	float64

## Dealing with outliers and empty values

In [9]: `data.fillna(data.median(), inplace= True)`

```
In [10]: max_length_aspect_ratio= data['max.length_aspect_ratio']
sn.boxplot(data= max_length_aspect_ratio)


for i in range(846):
    if max_length_aspect_ratio[i]> 13:
        max_length_aspect_ratio[i]= max_length_aspect_ratio.median()

data['max.length_aspect_ratio']= max_length_aspect_ratio
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel\_launcher.py:6: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))



```
In [11]: pr_access_aspect_ratio= data['pr.axis_aspect_ratio']
sn.boxplot(data= pr_access_aspect_ratio)


for i in range(846):
    if pr_access_aspect_ratio[i]> 80:
        pr_access_aspect_ratio[i]= pr_access_aspect_ratio.median()

data['pr.axis_aspect_ratio']= pr_access_aspect_ratio
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel\_launcher.py:6: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))




```
In [12]: radius_ratio= data['radius_ratio']
sn.boxplot(data= radius_ratio)

for i in range(846):
    if radius_ratio[i]> 250:
        radius_ratio[i]= radius_ratio.median()

data['radius_ratio']= radius_ratio
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))




```
In [15]: data.describe()
```

Out[15]:

	compactness	circularity	distance_circularity	radius_ratio	pr.axis_aspect_ratio	max.length_
count	846.000000	846.000000	846.000000	846.000000	846.000000	
mean	93.678487	44.823877	82.100473	168.230496	61.154846	
std	8.234474	6.134272	15.741569	32.018672	5.613458	
min	73.000000	33.000000	40.000000	104.000000	47.000000	
25%	87.000000	40.000000	70.000000	141.000000	57.000000	
50%	93.000000	44.000000	80.000000	167.000000	61.000000	
75%	100.000000	49.000000	98.000000	194.000000	65.000000	
max	119.000000	59.000000	112.000000	250.000000	76.000000	

```
In [16]: sn.pairplot(data, diag_kind = 'kde')
plt.show()
```



## Standardizing Data

```
In [17]: from sklearn.preprocessing import StandardScaler
data = StandardScaler().fit_transform(data)
```



## Splitting Data

```
In [18]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test= train_test_split(data, raw_data['class'], test_
```

## Implementing SVM

```
In [19]: from sklearn import svm
svma= svm.SVC()
svma.fit(x_train, y_train)
svm_prediction= svma.predict(x_test)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.  
 "avoid this warning.", FutureWarning)

```
In [20]: from sklearn.metrics import classification_report
from sklearn import metrics
print(classification_report(y_test,svm_prediction))
print("Accuracy:",metrics.accuracy_score(y_test, svm_prediction))
metrics.confusion_matrix(y_test, svm_prediction)
```

	precision	recall	f1-score	support
1	0.98	0.98	0.98	53
2	0.98	0.98	0.98	60
3	0.99	0.99	0.99	141
accuracy			0.99	254
macro avg	0.99	0.99	0.99	254
weighted avg	0.99	0.99	0.99	254

Accuracy: 0.9881889763779528

```
Out[20]: array([[ 52,   1,   0],
 [   0,  59,   1],
 [   1,   0, 140]], dtype=int64)
```

## K fold Validation

```
In [21]: from sklearn.model_selection import cross_val_score
clf = svm.SVC(kernel='linear', C=1)
scores = cross_val_score(clf, data, raw_data['class'], cv=5)
scores
```

```
Out[21]: array([0.94117647, 0.95294118, 0.94705882, 0.95857988, 0.95209581])
```



```
In [25]: svma= svm.SVC()
svma.fit(x_train, y_train)
svm_prediction= svma.predict(x_test)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.  
"avoid this warning.", FutureWarning)

```
In [26]: print(classification_report(y_test,svm_prediction))
print("Accuracy:",metrics.accuracy_score(y_test, svm_prediction))
metrics.confusion_matrix(y_test, svm_prediction)
```

	precision	recall	f1-score	support
1	0.92	0.92	0.92	53
2	0.97	0.98	0.98	60
3	0.98	0.97	0.98	141
accuracy			0.96	254
macro avg	0.96	0.96	0.96	254
weighted avg	0.96	0.96	0.96	254

Accuracy: 0.9645669291338582

```
Out[26]: array([[ 49,   2,   2],
 [   0,  59,   1],
 [   4,   0, 137]], dtype=int64)
```

## K fold validation

```
In [27]: from sklearn.model_selection import cross_val_score
clf = svm.SVC(kernel='linear', C=1)
scores = cross_val_score(clf, data, raw_data['class'], cv=5)
scores
```

```
Out[27]: array([0.94117647, 0.95294118, 0.94705882, 0.95857988, 0.95209581])
```

```
In [28]: print("Accuracy: %.2f (+/- %.2f)" % (scores.mean(), scores.std() * 2))
```

Accuracy: 0.95 (+/- 0.01)

```
In [30]: #To increase the accuracy, I have replaced outliers with medians and also standard
#In this case, applying PCA is reducing the accuracy. And in this case, PCA is not
#k fold validation for both the cases is same
#Without PCA:
#Accuracy: 98.8 kfold accuracy: 95
#With PCA:
#Accuracy: 96.4 kfold accuracy: 95
```

